

Neural Networks & Deep Learning - Saikumar Reddy Mandha(700765227)

GitHub Link: <https://github.com/saikumarreddyMandha/Assignment-7/tree/main>

Video Record : https://drive.google.com/file/d/1SWcUZEbOPzXle66drvvgTijO64AakAGC/view?usp=drive_link

Use Case Description:

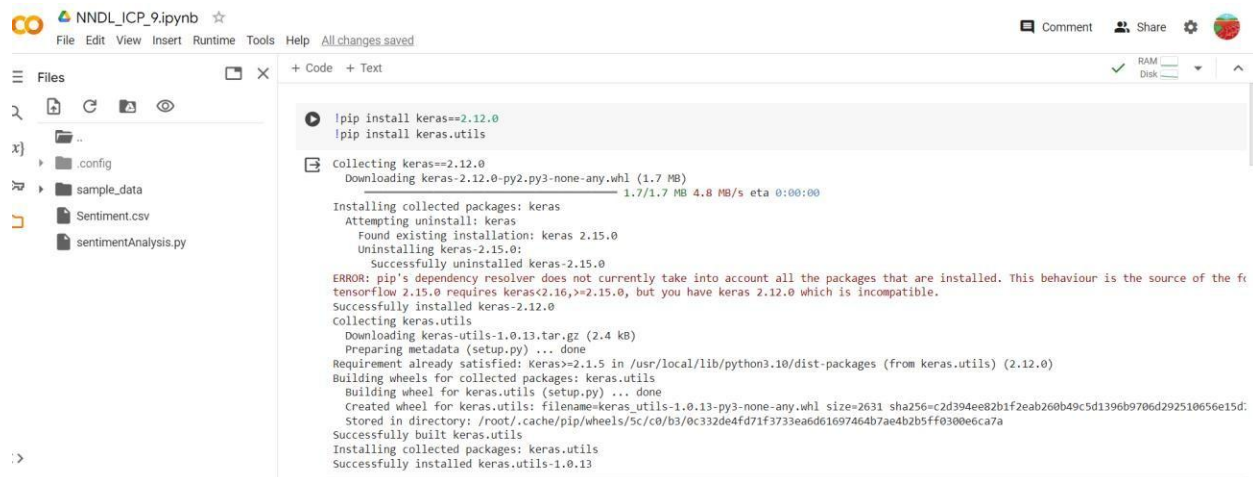
Sentiment Analysis on the Twitter dataset

Programming elements:

1. Basics of LSTM
2. Types of RNN
3. Use case: Sentiment Analysis on the Twitter data set

In class programming:

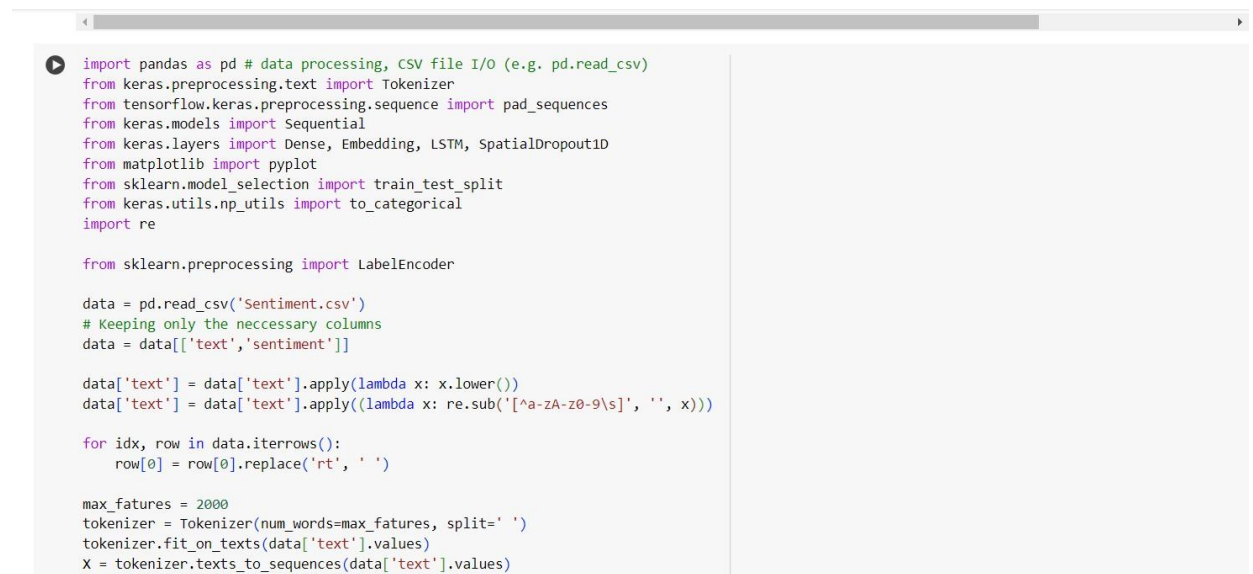
1. Save the model and use the saved model to predict on new text data (ex, "A lot of good things are happening. We are respected again throughout the world, and that's a great [thing.@realDonaldTrump](#)")
2. Apply GridSearchCV on the source code provided in the class.



The screenshot shows a Jupyter Notebook window titled "NNDL_ICP_9.ipynb". The left sidebar displays a file explorer with folders like ".config" and "sample_data", and files like "Sentiment.csv" and "sentimentAnalysis.py". The main area shows a code cell with the following output:

```
[pip install keras==2.12.0
pip install keras.utils

Collecting keras==2.12.0
  Downloading keras-2.12.0-py2.py3-none-any.whl (1.7 MB)
    1.7/1.7 MB 4.8 MB/s eta 0:00:00
Installing collected packages: keras
  Attempting uninstall: keras
    Found existing installation: keras 2.15.0
    Uninstalling keras-2.15.0:
      Successfully uninstalled keras-2.15.0
  Successfully installed keras-2.12.0
Collecting keras.utils
  Downloading keras-utils-1.0.13.tar.gz (2.4 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: Keras>=2.1.5 in /usr/local/lib/python3.10/dist-packages (from keras.utils) (2.12.0)
Building wheels for collected packages: keras.utils
  Building wheel for keras.utils (setup.py) ... done
  Created wheel for keras.utils: filename=keras_utils-1.0.13-py3-none-any.whl size=2631 sha256=c2d394ee82b1f2eab260b49c5d1396b9706d292510656e15d
  Successfully built keras.utils
Installing collected packages: keras.utils
Successfully installed keras.utils-1.0.13
```



The screenshot shows a Jupyter Notebook window with a code cell containing the following Python code:

```
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from keras.models import Sequential
from keras.layers import Dense, Embedding, LSTM, SpatialDropout1D
from matplotlib import pyplot
from sklearn.model_selection import train_test_split
from keras.utils.np_utils import to_categorical
import re

from sklearn.preprocessing import LabelEncoder

data = pd.read_csv('Sentiment.csv')
# Keeping only the necessary columns
data = data[['text', 'sentiment']]

data['text'] = data['text'].apply(lambda x: x.lower())
data['text'] = data['text'].apply(lambda x: re.sub('[^a-zA-Z0-9\s]', '', x))

for idx, row in data.iterrows():
    row[0] = row[0].replace('rt', ' ')

max_features = 2000
tokenizer = Tokenizer(num_words=max_features, split=' ')
tokenizer.fit_on_texts(data['text'].values)
X = tokenizer.texts_to_sequences(data['text'].values)
```

```

X = pad_sequences(X)

embed_dim = 128
lstm_out = 196
def createmodel():
    model = Sequential()
    model.add(Embedding(max_fatures, embed_dim,input_length = X.shape[1]))
    model.add(LSTM(lstm_out, dropout=0.2, recurrent_dropout=0.2))
    model.add(Dense(3,activation='softmax'))
    model.compile(loss = 'categorical_crossentropy', optimizer='adam',metrics = ['accuracy'])
    return model
# print(model.summary())

labelencoder = LabelEncoder()
integer_encoded = labelencoder.fit_transform(data['sentiment'])
y = to_categorical(integer_encoded)
X_train, X_test, Y_train, Y_test = train_test_split(X,y, test_size = 0.33, random_state = 42)

batch_size = 32
model = createmodel()
model.fit(X_train, Y_train, epochs = 1, batch_size=batch_size, verbose = 2)
score,acc = model.evaluate(X_test,Y_test,verbose=2,batch_size=batch_size)
print(score)
print(acc)
print(model.metrics_names)

```

```

291/291 - 55s - loss: 0.8254 - accuracy: 0.6419 - 55s/epoch - 191ms/step
144/144 - 3s - loss: 0.7654 - accuracy: 0.6660 - 3s/epoch - 24ms/step
0.7654296159744263
0.6660113334655762
['loss', 'accuracy']

```

```
[ ] model.save('sentiment_model.h5')
```

```

from keras.models import load_model
import numpy as np

loaded_model = load_model('sentiment_model.h5')

new_text = ["A lot of good things are happening. We are respected again throughout the world, and that's a great thing.@realDonaldTrump"]
new_text = tokenizer.texts_to_sequences(new_text)
new_text = pad_sequences(new_text, maxlen=X.shape[1], dtype='int32', value=0)
sentiment_prob = loaded_model.predict(new_text, batch_size=1, verbose=2)[0]

sentiment_classes = ['Positive', 'Neutral', 'Negative']
sentiment_pred = sentiment_classes[np.argmax(sentiment_prob)]

print("Predicted sentiment: ", sentiment_pred)
print("Predicted probabilities: ", sentiment_prob)

```

```

1/1 - 0s - 304ms/epoch - 304ms/step
Predicted sentiment: Positive
Predicted probabilities: [0.47510943 0.17564584 0.34924477]

```

This code loads the saved model using the `load_model` function, and then preprocesses the new text data in the same way as the training data. The `predict` method is called on the loaded model to get the predicted class probabilities for the new text data. The class with the highest probability is chosen as the predicted sentiment. The predicted sentiment and probabilities are then printed to the console.

To apply `GridSearchCV` on the provided source code, we can use the `GridSearchCV` class from `sklearn` to search for the best combination of hyperparameters for the LSTM model. The hyperparameters that can be tuned are the number of LSTM units, the dropout rate, and the learning rate of the optimizer.

```
from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
from keras.layers import LSTM

# Function to create the model, as it's required by KerasClassifier
def create_model(lstm_out=196, dropout=0.2):
    model = Sequential()
    model.add(Embedding(max_features, embed_dim, input_length=X.shape[1]))
    model.add(LSTM(lstm_out, dropout=dropout, recurrent_dropout=dropout))
    model.add(Dense(3, activation='softmax'))
    model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
    return model

# Create the KerasClassifier
model = KerasClassifier(build_fn=create_model, epochs=1, batch_size=batch_size, verbose=2)

# Define the grid of parameters to search
param_grid = {
    'lstm_out': [196, 256],
    'dropout': [0.2, 0.3]
}

# Create GridSearchCV
grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1, cv=3)
grid_result = grid.fit(X_train, Y_train)

# Summarize results
print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
```

<ipython-input-8-658eda5ed78a>:15: DeprecationWarning: KerasClassifier is deprecated, use Sci-Keras (<https://github.com/adriangb/scikeras>) instead
model = KerasClassifier(build_fn=create_model, epochs=1, batch_size=batch_size, verbose=2)
291/291 - 54s - loss: 0.8228 - accuracy: 0.6437 - 54s/epoch - 186ms/step
Best: 0.668568 using {'dropout': 0.2, 'lstm_out': 196}

This code defines the `create_model` function that returns a Keras model with the specified hyperparameters. The `KerasClassifier` class is used to create a wrapper for the `create_model` function, which can be used as an estimator for `GridSearchCV`. The hyperparameters to be tuned are defined in the `param_grid` dictionary. `GridSearchCV` is then called with the `KerasClassifier` object, the `param_grid` dictionary.